

DAY 1

ELK Pre-requisite:

Java

Python (Because ELK has a lot of rest services, you can write your own apis and in industry python is a way to go)

Software Requirements:

Windows 10/11 or Mac

16 GB RAM

ELK

Anaconda Navigator

Postman software

VS Code

Rancher Desktop

Docker Desktop

Monolithic Architecture:

Client

Server

Microservices Architecture:

We dont use any central servers.

Dont write everything under one umbrella.

Actuator framework for Microsrvice

ELK is not only used for logging, metrics but is also used for NLP (AI) for faster results etc.

ELK is HA architecture.

ELK stack

ELK: uses logstash

EFK: Uses Fluentd

Users can take any data from any source and perform searching, analyzing and creating visualizations.

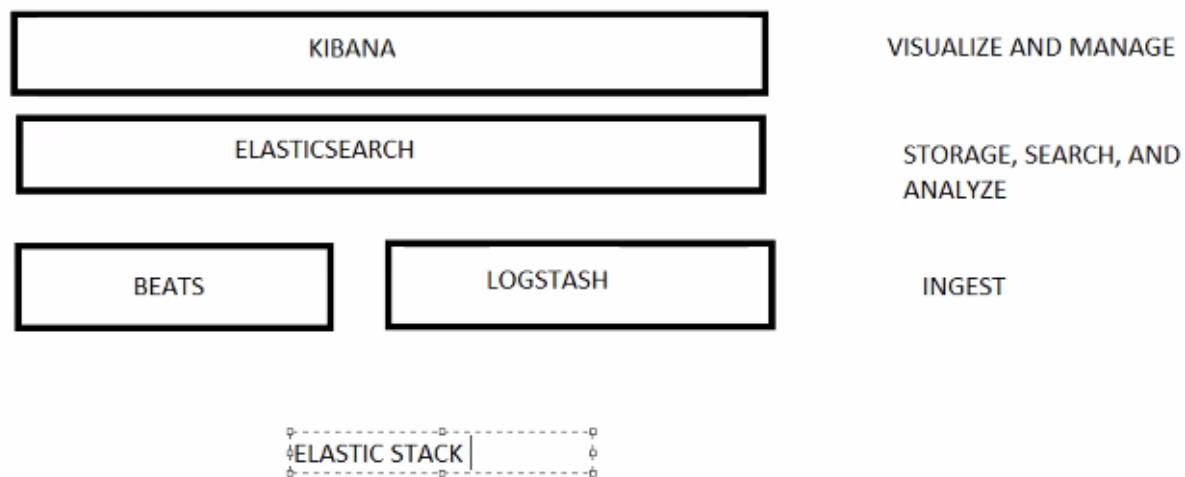
REST will help in understanding

Came in 2012 and from 2012 onwards people started using ELK

Components Of Elastic Stack:

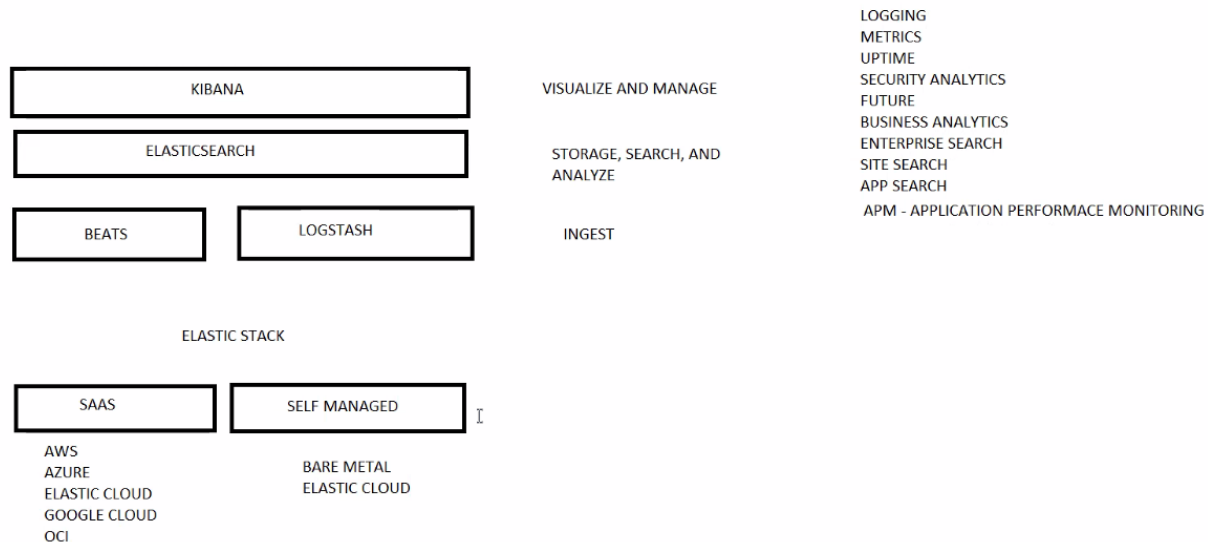
1. Elasticsearch: Store the data (SQL, NO SQL, JSON, CSV). Most important and basic component of ELK. It is open source and distributed and supports restful and JSON based.
2. Logstash: Can write data pipeline used to read the data from various sources and write data into different target data sources.
3. Kibana: GUI used to visualize the data (Analytics, dashboards)
4. Beats: Lightweight data shipper on every target machine and sends data to elastic search directly or via Logstash.

Apache Lucene: High performance text searching engine written in java.



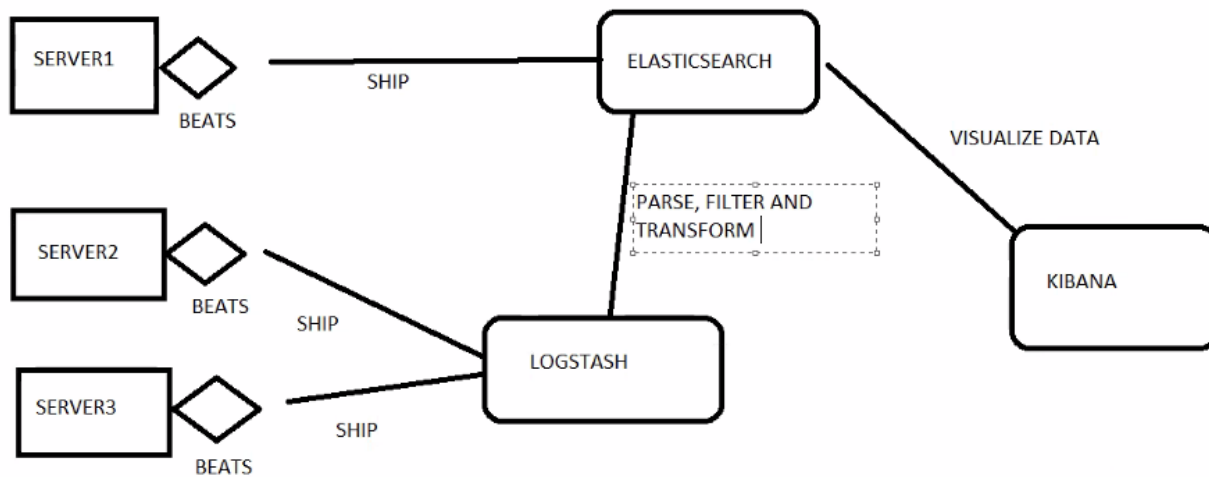
Services that ELK offers:

Logging, Metrics, Uptime, Security Analytics, Future, Business Analytics, Enterprise Search, Site search, App search, APM



Two ways to send data:

1. Server sending directly to Elastic Search and display on Kibana
2. Server sends data to elastic search via Logstash and visualize data on kibana.



We can write our own API for our data pipeline flow, no need for LogStash, just send it directly to Elastic Search. Logstash parses, filters and transforms the data and sends it but if you already have something that can parse, filter or transform the data then Logstash can be skipped.

ElasticSearch:

Full text search engine. NO SQL DB. Can be used as analytics engine. Schema less and works in real time. Written in java rest APIs. 12 factor app design. Structured and unstructured data can be used

Kibana:

It's a GUI used to do multiple activities. It discovers your data by exploring. It can analyze the data and implement the metrics. It can visualize the data. It can apply ML on the data (forecasting, anomalies, etc). It can do APM. It can Manage roles and users. Elastic search expressions. It can analyze time series data using the timeline.

LogStash:

It's a data pipeline. Push the data into Elasticsearch. You can take any type of data using Logstash. It parses, filters and transforms the data.

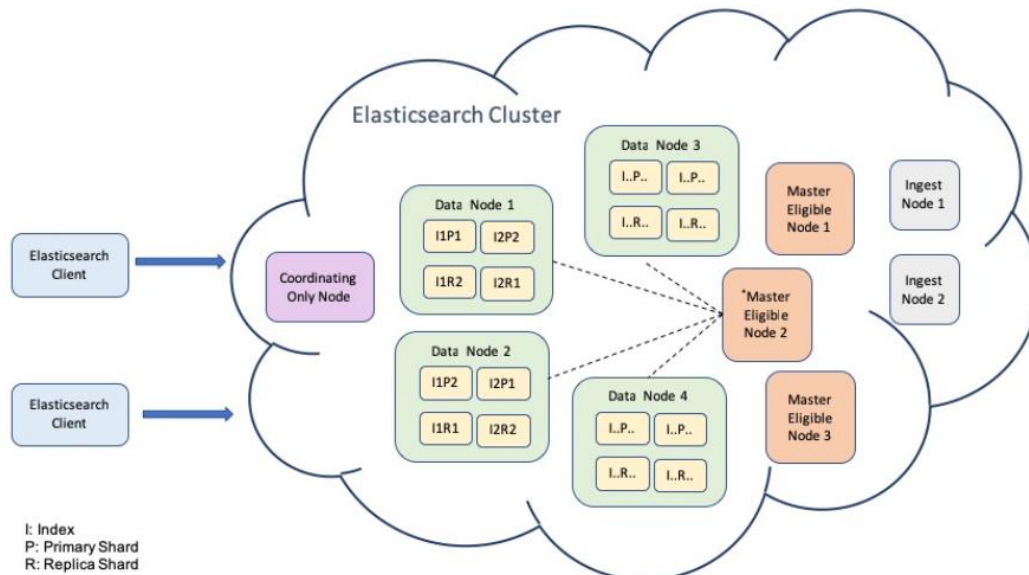
Beats:

It's a lightweight software agent. It can talk to logstash or Elasticsearch directly. There are different beats such as for network, windows, etc we have different beats agent (file beat, winlog beat packet beat).

X-pack:

It is a security extension. Elastic stack extension to provide security. It provides secured Alerting, secured monitoring, secured reporting, and secured ML. It is not free.

An Elasticsearch cluster is a group of one or more Elasticsearch nodes that are connected together. Let's first outline how it is laid out, as shown in the following diagram:



Data node: Contains the data that contains indexed documents. Performs CRUD operation.

Ingest Node: One of the ways to process the document in your pipeline mode before indexing the document.

Disable it using `node.ingest=false` in `elasticsearch.yml` file

Coordinating Only Node: Disable data node ingest node. At such time they can only route the request and distribute work using a bulk indexing mechanism.

Disabling the different components will tell elastic search what role is it playing.

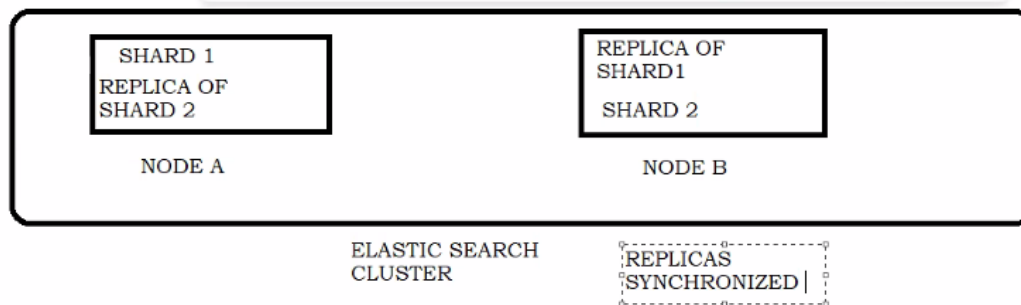
Fields: Smallest individual unit of data.

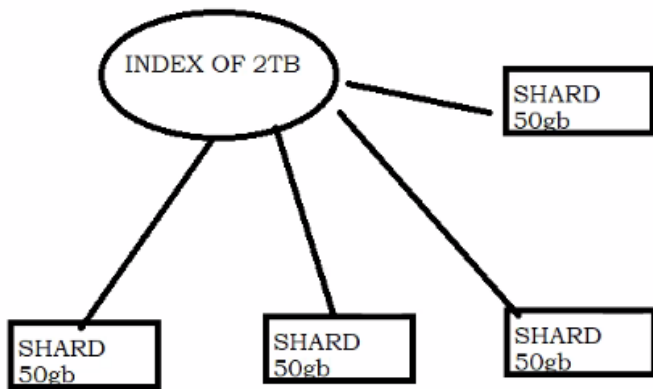
Multifields: Can be indexed in more than one way to produce more results.

Shards: Single Lucene index

Index size is a major cause of elastic search collapse. Indexing can begin to fail. Therefore, shards are heavily used. Split the indexes horizontally into pieces. Helps in distributed environment and will have multiple nodes.

Replica: Copies of indexes. Replicas are not placed on the same nodes as shards.





Sharding: Improves performance, Provides high availability and fault tolerance.

Installation:

Download the zip packages for all three components from here:

<https://www.elastic.co/downloads/elasticsearch>

<https://www.elastic.co/downloads/logstash>

<https://www.elastic.co/downloads/kibana>

Now create a folder ELK and put all three zip files under this and extract all three.

Install jdk if not already done.

Start the elastic search using below:

C:\Users\abratuiri\Desktop\ELK\elasticsearch-8.8.0\bin\elasticsearch.bat

Edit the below file to disable the security feature: xpack.security.enabled: false

C:\Users\abratuiri\Desktop\ELK\elasticsearch-8.8.0\config\elasticsearch.yml

Now restart the Elasticsearch

Hit the URL in your browser to see the information (if prompts for credentials appear just cancel it):

<http://127.0.0.1:9200/>

```

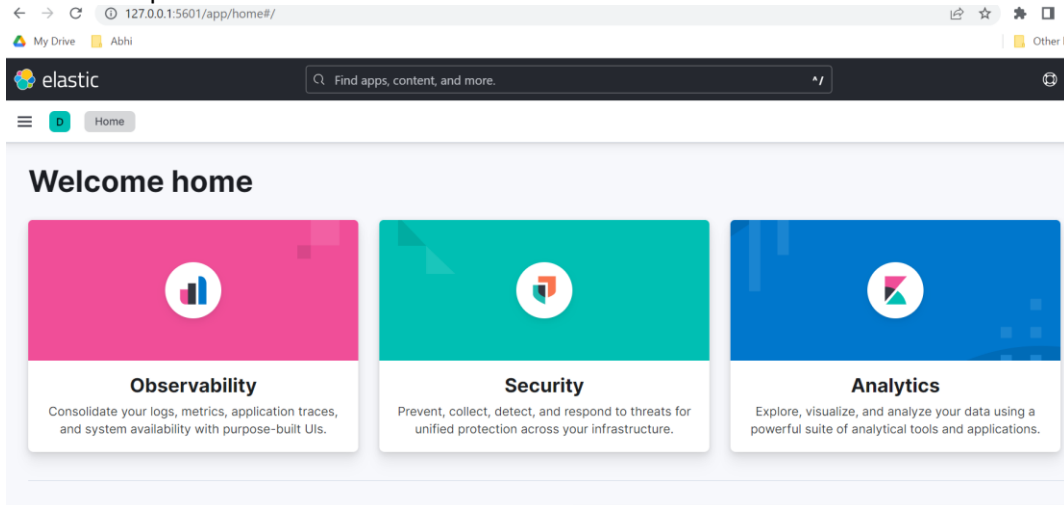
{
  name: "ABRATUR-I-X40-J",
  cluster_name: "elasticsearch",
  cluster_uuid: "15rOdpN0RrCwmKGjvPkPnQ",
  - version: {
    number: "8.8.0",
    build_flavor: "default",
    build_type: "zip",
    build_hash: "c01029875a091076ed42cdb3a41c10b1a9a5a20f",
    build_date: "2023-05-23T17:16:07.179039820Z",
    build_snapshot: false,
    lucene_version: "9.6.0",
    minimum_wire_compatibility_version: "7.17.0",
    minimum_index_compatibility_version: "7.0.0"
  },
  tagline: "You Know, for Search"
}
  
```

You can also use Postman to send the request and you will see the json results similar to what we saw in the browser.

Now start Kibana:

C:\Users\abraturi\Desktop\ELK\kibana-8.8.0\bin\kibana.bat

The default port for Kibana is 5601



Now create a file logstash.conf under path: C:\Users\abraturi\Desktop\ELK\logstash-8.8.0\logstash.conf

Logstash.conf:

```
input {
  stdin {}
}

output {
  elasticsearch {
    hosts => ["localhost:9200"]
    index => "indexforlogstash"
  }
}
```

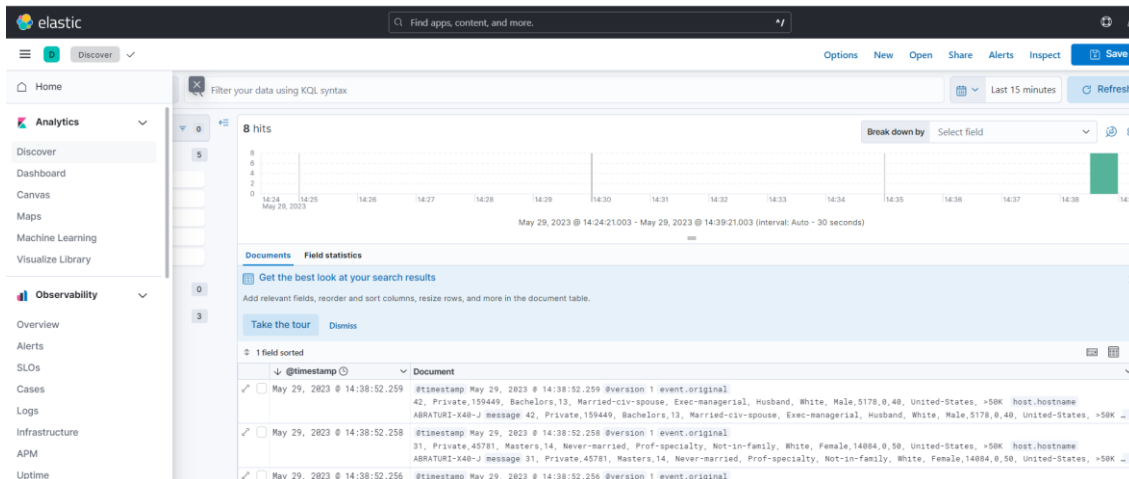
Now go to directory: C:\Users\abraturi\Desktop\ELK\logstash-8.8.0\bin

Once done execute the below command: `.\logstash.bat -f .\logstash.conf`

Now in this same terminal or cmd provide the data by just copying and pasting.

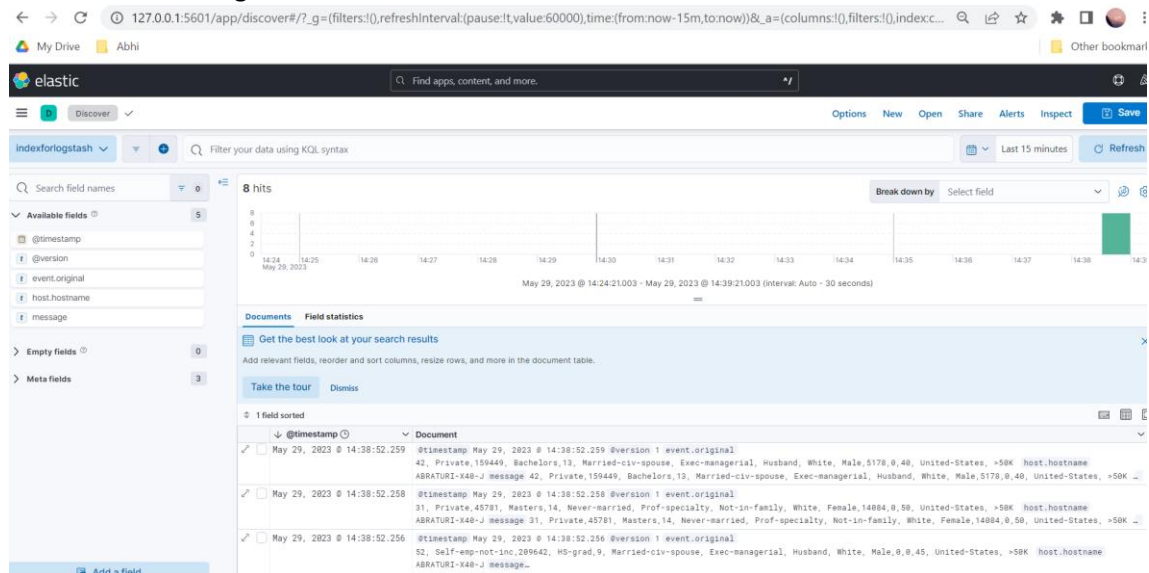
Check the data in Kibana:

Go to discover:



For the first time provide the search name and indexname ("indexforlogstash" in this case) and click on create search.

You will start seeing the data:



If you want to change the data path and log path then update elasticsearch.yaml with the below and restart elasticsearch server:

#admin defined configuration changes

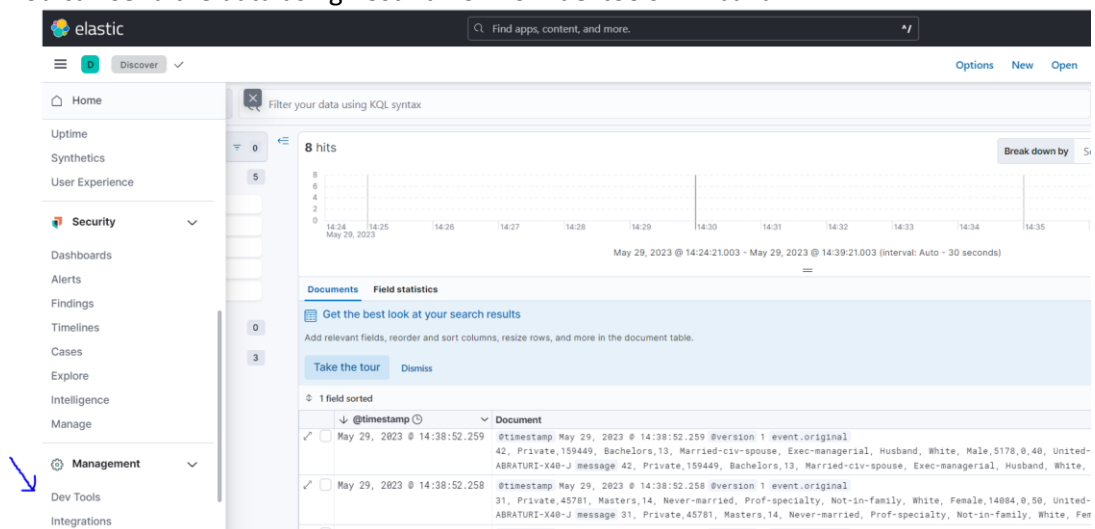
path:

```
data : "C:\\Users\\abratur1\\Desktop\\ELK\\elasticsearch-8.8.0\\data"
logs : "C:\\Users\\abratur1\\Desktop\\ELK\\elasticsearch-8.8.0\\logs"
```

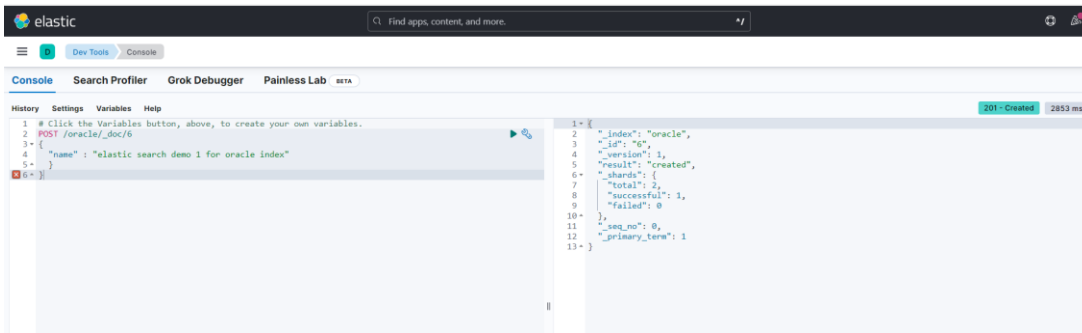
Start the elasticsearch:

C:\\Users\\abratur1\\Desktop\\ELK\\elasticsearch-8.8.0\\bin\\elasticsearch

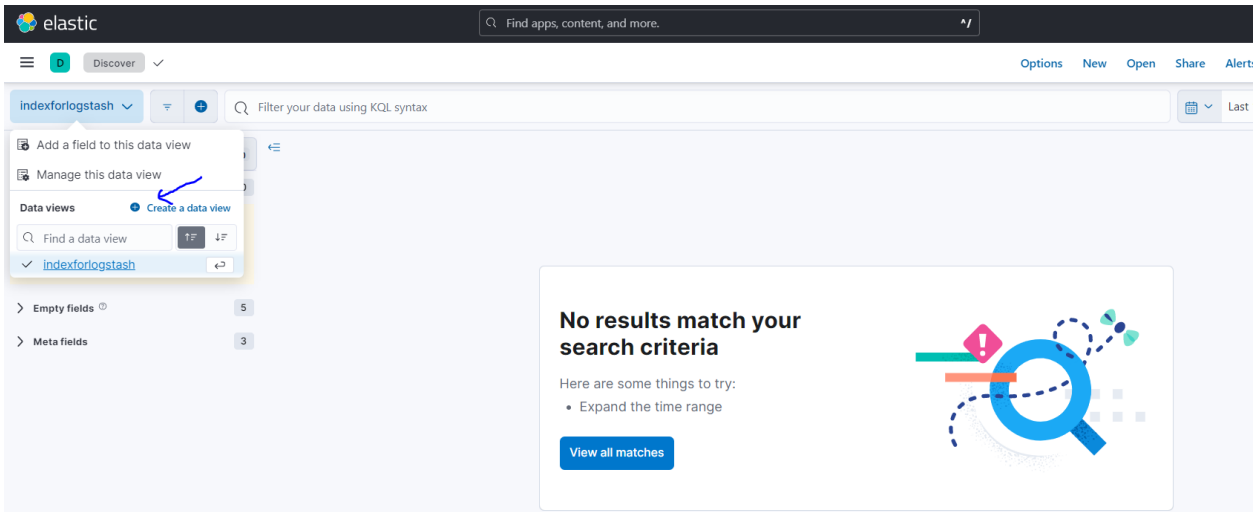
You can send the data using Postman or from devtools in kibana.



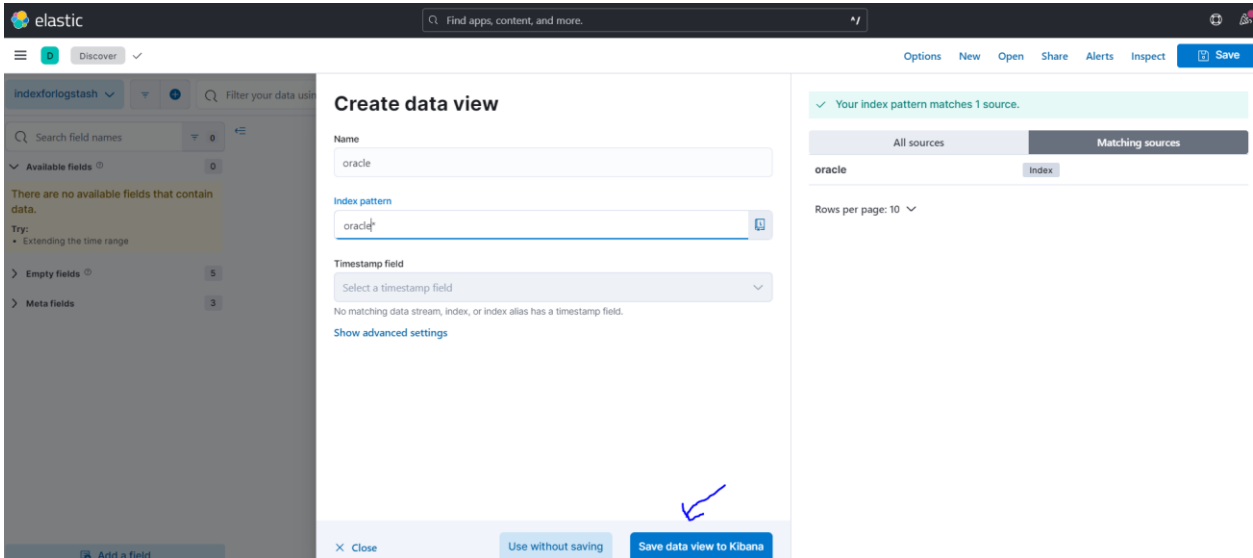
Now, post the query and run



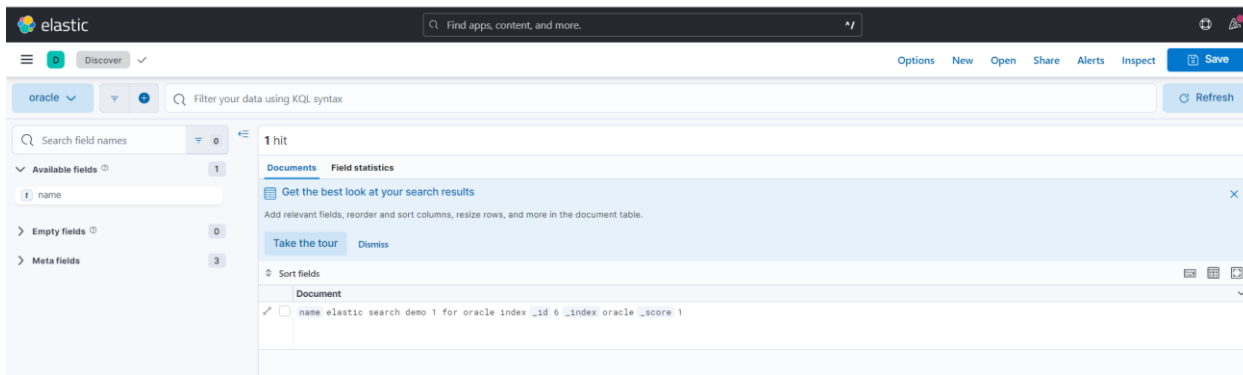
Now to check the logs sent via the POST method, create another data view that will see the logs in the index “oracle”



Fill the details and click on “Save data view”



You will see the data:



Performing Operations using curl command:

Create index:

curl --location --request PUT "http://localhost:9200/cf_etf"

Status of index:

curl --location --head http://localhost:9200/cf_etf

Get details of an index:

curl --location --request GET http://localhost:9200/cf_etf

Get details of an index:

curl --location --request GET <http://localhost:9200/indexforlogstash>

Delete an index:

curl --location --request DELETE http://localhost:9200/cf_etf

Close the index and it will not show in kibana (GUI)

curl --location --request POST http://localhost:9200/cf_etf/_close

Open the index and it show in kibana (GUI)

curl --location --request POST http://localhost:9200/customer/_open

You can also change the content by REST calls:

curl --location --request PUT "http://localhost:9200/customer/_settings" --header "Content-Type: application/json" --data '{"index\" : { \"number_of_replicas\" : 5 } }'

curl --location --request GET http://localhost:9200/customer/_settings

Creating index template using Postman (REST calls):

Create customer index template

...

http://localhost:9200/_template/customer_template

PUT

Header - content-type - application/json

BODY

```
{
  "index_patterns" : ["customer*"],
  "settings" : {
    "number_of_replicas" : 3,
    "codec" : "best_compression"
  }
}
```

```
}  
...
```

get customer index template

```
...
```

http://localhost:9200/_template/customer_template

GET

Header - content-type - application/json

```
...
```

check if exists customer index template

```
...
```

http://localhost:9200/_template/customer_template

HEAD

Header - content-type - application/json

```
...
```

check if exists customer index template

```
...
```

http://localhost:9200/_template/customer_template

DELETE

Header - content-type - application/json

```
...
```

DAY 2

Alias points to the main object itself. It is used to do pointing capability. You cannot do reading and writing simultaneously in elasticsearch (cannot do on a single object).

Creating an index:

```
...
```

<http://localhost:9200/oracle1>

PUT

Header - content-type - application/json

```
...
```

Creating alias for index "oracle":

```
...
```

http://localhost:9200/oracle/_alias/oracle_alias

PUT

Header - content-type - application/json

```
...
```

Change alias from one index to another index with zero downtime:

```
...
```

http://localhost:9200/_aliases

POST

Header - content-type - application/json

```
...
...
BODY
{
  "actions" : [
    {"remove" : {"index" : "cf_etf", "alias" : "cf_etf_production"}},
    {"add" : {"index" : "cf_etf_new", "alias" : "cf_etf_production"}}
  ]
}
...
```

Get the alias details for all indexes:

```
...
http://localhost:9200/oracle/\_alias
GET
...
```

Get the alias details of all indexes that has alias "oracle__alias":

```
...
http://localhost:9200/\*/ \_alias/oracle\_\_alias
GET
...
```

Creating a view:

```
...
http://localhost:9200/cf\_view
PUT
...
```

```
...
BODY
{
  "mappings" : {
    "properties" : {
      "symbol" : {"type" : "keyword"},
      "category" : {"type" : "keyword"}
    }
  }
}
...
```

Creating alias with filter:

```
...
http://localhost:9200/\_aliases

POST
...
...
BODY
{
  "actions" : [
    {"add" : {"index" : "cf_view", "alias" : "cf_view_international", "filter" : {"term" :
{"category" : "International"}}}}
  ]
}
```

```
    ]  
  }  
  ...
```

Get the search details from an index:

...

<http://localhost:9200/cf view international/ search>

GET

...

Get the full statistics details for an index:

...

<http://localhost:9200/cf view/ stats>

GET

...

Get the segment details for an index:

...

<http://localhost:9200/cf view/ segments>

GET

...

Get the shard details for an index:

...

<http://localhost:9200/cf view/ shard stores>

GET

...

Changing the settings of index:

...

<http://localhost:9200/cf etf/ settings>

PUT

...

...

BODY

```
{  
  "settings" : {  
    "index.blocks.write" : true  
  }  
}  
...
```

Create the shards of an index:

...

<http://localhost:9200/cf etf/ split/cf etf split>

POST

...

```
{  
  "settings" : {  
    "index.number_of_shards" : 4  
  }  
}
```

```
}
```

Creating the document and sending the data (Using Postman):

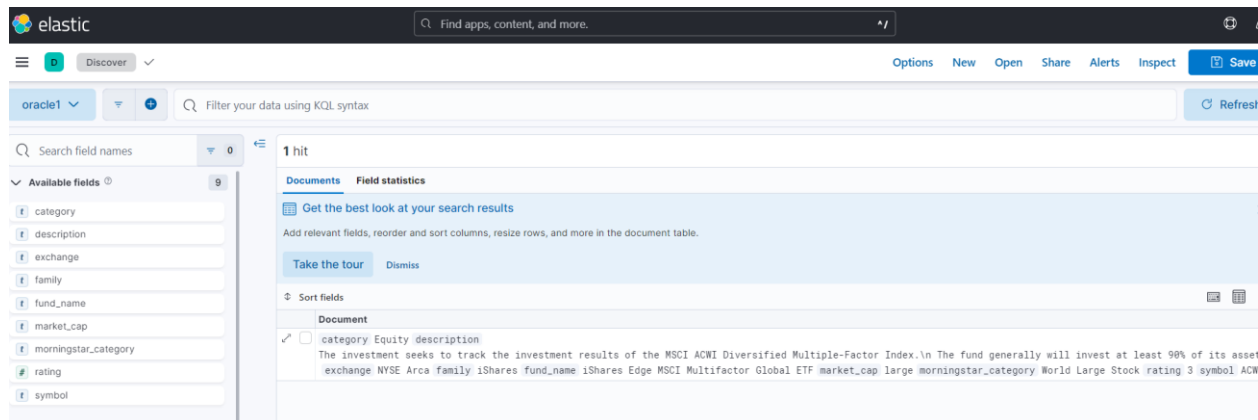
```
...
```

<http://localhost:9200/cf etf/ doc/1>

PUT

```
...
```

```
{
  "symbol": "ACWF",
  "fund_name": "iShares Edge MSCI Multifactor Global ETF",
  "rating": 3,
  "morningstar_category": "World Large Stock",
  "category": "Equity",
  "family": "iShares",
  "market_cap": "large",
  "description": "The investment seeks to track the investment results of the MSCI ACWI Diversi
fied Multiple-
Factor Index.\n\n The fund generally will invest at least 90% of its assets in the component se
curities of the underlying index and in investments that have economic characteristics that ar
e substantially identical to the component securities of the underlying index. The underlying
index is designed to contain equity securities from the MSCI ACWI Index (the \\\\\"parent index\\
\\") that have high exposure to four investment style factors: value, quality, momentum and lo
w size, while maintaining a level of risk similar to that of the parent index.",
  "exchange": "NYSE Arca"
}
```



Getting the selected fields only from the data using the partial search:

```
...
```

<http://localhost:9200/oracle1/ doc/1? source=symbol.category.rating>

GET

```
...
```

Overview GET http: PUT http: POST http: PUT http: POST http: GET http: + ... No Environment

http://localhost:9200/cf_etf/_split/cf_etf_split

GET http://localhost:9200/oracle1/_doc/?_source=symbol,category,rating

Params Authorization Headers (8) Body Pre-request Script Tests Settings Cookies

Query Params

Key	Value	Description	...	Bulk Edit
<input checked="" type="checkbox"/> _source	symbol,category,rating			
Key	Value	Description		

Body Cookies Headers (4) Test Results 200 OK 8 ms 273 B Save as Example

Pretty Raw Preview Visualize JSON

```
7  {
8    "found": true,
9    "_source": {
10      "symbol": "ACWF",
11      "rating": 3,
12      "category": "Equity"
13    }
14  }
```

Update the selected fields only in the data:

...

http://localhost:9200/oracle1/_update/1

POST

...

...

BODY

```
{
  "doc" : {
    "market_cap" : "mid"
  }
}
```

...

Multiple Documents:

Get multiple docs:

...

http://localhost:9200/_mget

POST

...

...

BODY

```
{
  "docs" : [
    { "_index" : "cf_view", "_id" : "1", "_source" : ["symbol","category"]}
  ]
}
```

...

Get docs by id:

...

http://localhost:9200/_mget

POST

...

...

BODY

```
{
  "docs" : [
    { "_index" : "cf_view", "_id" : "1", "_source" : ["symbol","category"] },
    { "_index" : "cf_etf", "_id" : "1", "_source" : ["symbol","category"] }
  ]
}
```

...

Perform bulk actions:

...

http://localhost:9200/_bulk

POST

...

...

BODY

```
{"delete": {"_index": "cf_etf", "_id": "1"}}
{"create": {"_index": "cf_etf", "_id": "1"}}
{"symbol": "ACWF", "fund_name": "iShares Edge MSCI Multifactor Global ETF", "rating": 3,
"morningstar_category": "World Large Stock", "category": "Equity", "family": "iShares",
"market_cap": "large", "description": "The investment seeks to track the investment results of
the MSCI ACWI Diversified Multiple-Factor Index.\n\nThe fund generally will invest at least
90% of its assets in the component securities of the underlying index and in investments that
have economic characteristics that are substantially identical to the component securities of
the underlying index. The underlying index is designed to contain equity securities from the
MSCI ACWI Index (the \"parent index\") that have high exposure to four investment style
factors: value, quality, momentum and low size, while maintaining a level of risk similar to
that of the parent index.\", \"exchange\": \"NYSE Arca\"}
```

Update by Query:

...

http://localhost:9200/_ctf/_update_by_query

POST

...

...

BODY

```
{"script": "if(ctx._source['rating']==3) {ctx._source['rating']==4}"}
```

Reindex:

...

http://localhost:9200/_reindex

POST

...

...

BODY

```
{
  "source": {"index": "cf_view"},
  "dest": {"index": "cf_etf"}
}
```

Mapping:

Dynamic mapping: Based on type or values maps will be created

Custom Mapping: User explicitly maps the content

Create default mapping index:

...

http://localhost:9200/default_mappins_index/

PUT

...

...

BODY

```
{
  "mappings" : {
    "dynamic_date_formats" : ["yyyy/MM/dd"]
  }
}
```

Creating doc with default mapping index:

...

http://localhost:9200/default_mappins_index/doc/1

PUT

...

...

BODY

```
{
  "simple_string_value": "a string",
  "string_w_date_value": "2019/02/07",
  "string_w_floating_point_value": "1.11",
  "string_w_integer_value" : "1",
  "integer_value": 1,
  "floating_pionnt_value": 1.1,
  "json_object" : {"level_1" : {"level2_1": "a string", "level_2_2": 1,
"level_2_3": 1.11}},
  "array_of_integer": [1, 2],
  "array_of_float": [1.11, 2.22],
  "array_of_simple_string": ["ab", "cd"],
  "array_of_integer_string": ["1", "2"],
  "array_of_float_string": ["1.11", "2.22"],
  "boolean_value" : true,
  "null_value": null
}
```

Search the default mapping index:

...

http://localhost:9200/default_mappins_index/search

GET

...

Create the custom mapping index with different structure:

...

http://localhost:9200/custom_mappins_index/

PUT

...

...

BODY

```
{
  "mappings" : {
    "date_detection" : false,
    "numeric_detection" : true
  }
}
```

Get the custom mapping index details:

...

http://localhost:9200/custom_mappings_index/_mappings

PUT

...

Perform the join in Dev tools in Kibana:

PUT test

PUT test/_mapping

```
{
  "properties" : {
    "id" : {"type" : "keyword"},
    "date" : {"type" : "date"},
    "customer_id" : {"type" : "keyword"},
    "sent" : {"type" : "boolean"},
    "item" : {"type" : "object",
      "properties" : {
        "name" : {"type" : "keyword"},
        "quantity" : {"type" : "long"},
        "cost" : {"type": "alias", "path": "item.price"},
        "price" : {"type" : "double"},
        "vat" : {"type" : "double"}
      }
    }
  }
}
```

PUT test/_doc/1?refresh

```
{
  "id": "1",
  "date": "2018-11-16T20:07:45Z",
  "customer_id": "100",
  "sent": true,
  "item": [
    {
      "name": "tshirt",
      "quantity": 10,
      "price": 4.3,
      "vat": 8.5
    }
  ]
}
```

```
}
```

GET test/_mapping

GET test/_search

```
{
  "query": {
    "term": {
      "item.cost" : 4.3
    }
  }
}
```

Q. Why there is a need to create one alias for multiple indexes? Tell one scenario or a case study on this?

Uber has a marketplace data real time data analysis with elasticsearch. Uber marketplace app is a supply and demand. Interaction between them is creating trips and fare.

DAY 3

Analyzer:

Mainly used for doing analysis. Helps you to do text analysis (unstructured data). It can do data mining, some amount of linguistic search. Performs full text searching where it can search relevant results rather than exact search. Tokenization, normalization (lemmatization, stemming) is used.

Analyzer does not have NLP model but has a similar logic behind the analyzer. But you can use NLP model as an analyzer as well (custom setup).

Elasticsearch used standard analyzer.

Different analyzers are there e.g. whitespace analyzer, etc.

Creating Analyzer:

...

http://localhost:9200/_analyze

POST

...

HEADER

Content-type application/json

...

BODY

```
{
  "text" : "<body> You all working on Elastic search </body>",
  "tokenizer" : "standard"
}
```

Creating lowercase Analyzer:

...

http://localhost:9200/_analyze

POST

...

HEADER

Content-type application/json

...

BODY

```
{
```

```
    "text" : "<body> You'll working on Elastic search </body>",
    "tokenizer" : "standard",
    "filter" : ["lowercase"]
}
```

Creating Analyzer with filter:

...

http://localhost:9200/_analyze

POST

...

HEADER

Content-type application/json

...

BODY

```
{
    "text" : "<body> You'll working on Elastic search </body>",
    "char_filter" : ["html_strip", {"type" : "mapping", "mappings" : ["You'll => You will"]}],
    "tokenizer" : "standard",
    "filter" : ["lowercase"]
}
```

Using Analyzer to replace:

...

http://localhost:9200/_analyze

POST

...

HEADER

Content-type application/json

...

BODY

```
{
    "text" : "<body> You'll working on Elastic search 8.8</body>",
    "char_filter" : ["html_strip", {"type" : "mapping", "mappings" : ["You'll => You will"]}],
    {"type" : "pattern_replace", "pattern" : "(\\d+).(\\d+)", "replacement" : "v$1"}],
    "tokenizer" : "standard",
    "filter" : ["lowercase"]
}
```

Using Analyzer for URL:

...

http://localhost:9200/_analyze

POST

...

HEADER

Content-type application/json

...

BODY

```
{
    "text" : "POST https://api.iextrading.com/1.0/stock/acwf/company /usr/local",
    "tokenizer" : "standard"
}
```

Using whitespace Analyzer for URL:

...

<http://localhost:9200/analyze>

POST

...

HEADER

Content-type application/json

...

BODY

```
{
  "text" : "POST https://api.iextrding.com/1.0/stock/acwf/company /usr/local",
  "tokenizer" : "whitespace"
}
```

Using uax_url_email Analyzer for URL (URL is retained):

...

<http://localhost:9200/analyze>

POST

...

HEADER

Content-type application/json

...

BODY

```
{
  "text" : "POST https://api.iextrding.com/1.0/stock/acwf/company /usr/local",
  "tokenizer" : "uax_url_email"
}
```

Ngram tokenizer in Analyzer:

...

<http://localhost:9200/analyze>

POST

...

HEADER

Content-type application/json

...

BODY

```
{
  "text" : "ElasticSearch 8.8",
  "tokenizer" : { "type" : "ngram", "min-
gram" : 2, "max_gram" : 2, "token_chars" : ["punctuation", "digit"] }
}
```

Simple_pattern tokenizer in Analyzer:

...

<http://localhost:9200/analyze>

POST

...

HEADER

Content-type application/json

...

BODY

```
{
  "text" : "ElasticSearch 8.8",
  "tokenizer" : { "type" : "simple_pattern", "pattern" : "[a-zA-Z]*" }
}
```

http://localhost:9200/_analyze

POST

HEADER

Content-type application/json

BODY

```
{
  "text" : "<body> You'll working on Elastic search 8.8</body>",
  "tokenizer" : "standard",
  "filter" : ["fingerprint"]
}
```

Analyzer with stemming filter:

http://localhost:9200/_analyze

POST

HEADER

Content-type application/json

BODY

```
{
  "text" : "love loves loved",
  "tokenizer" : "standard",
  "filter" : [{"type": "stemmer", "name" : "english"}]
}
```

Analyzer with stemming filter returning only unique stemming word:

http://localhost:9200/_analyze

POST

HEADER

Content-type application/json

BODY

```
{
  "text" : "love loves loved",
  "tokenizer" : "standard",
  "filter" : [{"type": "stemmer", "name" : "english"}, "unique"]
}
```

Analyzer with filter and script:

...

http://localhost:9200/_analyze

POST

...

HEADER

Content-type application/json

...

BODY

```
{
  "text" : "Elasticsearch 8.8 is the search and analytics engine that powers the Elastic Engine",
  "tokenizer" : "standard",
  "filter" : [{"type": "condition", "script" : {"source" : "token.getType()=='<ALPHANUM>'"},
  "filter" : ["reverse"]}]]
}
```

Analyzer with delimiter:

...

http://localhost:9200/_analyze

POST

...

HEADER

Content-type application/json

...

BODY

```
{
  "text" : "Elasticsearch 8.8 is the search and analytics engine that powers the Elastic Engine",
  "tokenizer" : "standard",
  "filter" : [{"type": "word_delimiter", "generate_word_parts" : true }]
}
```

Analyzer with catenate:

...

http://localhost:9200/_analyze

POST

...

HEADER

Content-type application/json

...

BODY

```
{
  "text" : "Elasticsearch 8.8 is the search and analytics engine that powers the Elastic Engine",
  "tokenizer" : "standard",
  "filter" : [{"type": "word_delimiter", "catenate_all" : true }]
}
```

Analyzer with split_on_case_change:

```
...  
http://localhost:9200/\_analyze  
POST  
...  
HEADER  
Content-type application/json  
...  
BODY  
{  
  "text" : "Elasticsearch 8.8 is the search and analytics engine that powers the Elastic Eng  
ine",  
  "tokenizer" : "standard",  
  "filter" : [{"type": "word_delimiter", "split_on_case_change" : true }]  
}
```

Standard Analyzer:

```
...  
http://localhost:9200/\_analyze  
POST  
...  
HEADER  
Content-type application/json  
...  
BODY  
{  
  "text" : "In Elasticsearch 8.8",  
  "analyzer" : "standard"  
}
```

Whitespace Analyzer:

```
...  
http://localhost:9200/\_analyze  
POST  
...  
HEADER  
Content-type application/json  
...  
BODY  
{  
  "text" : "In Elasticsearch 8.8",  
  "analyzer" : "whitespace"  
}
```

Other Analyzers: fingerprint, keyword, English, pattern, simple

NOTE: Analyzer is a wrapper on top of entities whereas tokenizer is used with filters. We can create custom analyzer under the index.

Custom Analyzer:

First create file "stopwords" under the config in elasticsearch.

...

http://localhost:9200/cf_etf

PUT

...

HEADER

Content-type application/json

...

BODY

```
{
  "mappings": {
    "dynamic": false,
    "properties": {
      "symbol": {"type": "keyword"},
      "fund_name": {"type": "text", "fields": {"raw": {"type": "keyword"}}},
      "rating": {"type": "byte"},
      "morningstar_category": {"type": "keyword"},
      "category": {"type": "keyword"},
      "family": {"type": "keyword"},
      "market_cap": {"type": "keyword"},
      "description": {"type": "text", "fields": {"raw": {"type": "keyword"}}},
      "exchange": {"type": "keyword"}
    }
  },
  "settings": {
    "analysis": {
      "analyzer": {
        "description_analyzer": {
          "type": "custom",
          "tokenizer": "description_tokenizer",
          "filter": ["lowercase", "description_pattern_replace_filter",
            "description_stemmer_filter", "description_stop_filter", "description_length_filter",
            "unique"]
        },
        "tokenizer": {
          "description_tokenizer": {"type": "char_group",
            "tokenize_on_chars": ["whitespace", "digit", "symbol", "\\n", ",", ":", "!", "?", ";", " ", "_",
              "{", "[", "]", "(", ")", "\\", "\\/", "\\\""]}
        },
        "filter": {
          "description_pattern_replace_filter": {"type": "pattern_replace",
            "pattern": "(\\w{2,})\\.", "replacement": "$1"},
          "description_stemmer_filter": {"type": "stemmer",
            "name": "light_english"},
          "description_stop_filter": {"type": "stop",
            "stopwords_path": "stopwords"},
          "description_length_filter": {"type": "length", "min": 2}
        }
      }
    }
  }
}
```



```
}
```

Now analyze the text using the above custom analyzer that was created:

```
...
```

http://localhost:9200/cf_etf/analyze

GET

```
...
```

HEADER

Content-type application/json

```
...
```

BODY

```
{
```

```
    "text": "The investment seeks to track the price and yield performance, before fees and expenses, of the Bloomberg Barclays U.S. Aggregate Enhanced Yield Index (the \\\\"index\\\\").\\n\\n The index is designed to broadly capture the U.S. investment grade, fixed income securities market while seeking to enhance yield within desired risk parameters and constraints. The fund is non-diversified.",
```

```
    "analyzer": "description_analyzer"
```

```
}
```

Now using analyzer with search:

git clone https://github.com/vishymails/ELK_RAW.git

```
curl -XPUT localhost:9200/cf_etf -H "Content-Type:application/json" --data-binary @define_custom_analyzer.json
```

```
curl -XPOST localhost:9200/cf_etf/_bulk?pretty -H "Content-Type:application/json" --data-binary @cf_etf_list_bulk.json
```

Now check the results using REST calls in Postman:

```
...
```

http://localhost:9200/cf_etf/search?from=2&size=2&sort=symbol:asc&source=false

POST

```
...
```

HEADER

Content-type application/json

```
...
```

```
...
```

http://localhost:9200/cf_etf/search?q=fund_name:ishares edge gobal&from=2&size=2&sort=score:asc

POST

```
...
```

HEADER

Content-type application/json

...

...

http://localhost:9200/cf_etf/_search?q=fund_name:ishares%20edge%20global&from=2&default_operator=AND&analyzer=keyword

POST

...

HEADER

Content-type application/json

...

...

http://localhost:9200/cf_etf/_search?q=fund_name:ishares%20edge%20global&from=2&default_operator=AND&explain=true

POST

...

HEADER

Content-type application/json

...

You can also search and filter:

...

http://localhost:9200/cf_etf/_search

POST

...

HEADER

Content-type application/json

...

BODY

```
{
  "query" : {
    "query_string" : {
      "query" : "family:iShares",
      "default_operator" : "AND"
    }
  }
}
```

...

http://localhost:9200/cf_etf/_search

POST

...

HEADER

Content-type application/json

...

BODY

```

{
  "query" : {
    "query_string" : {
      "query" : "family:iShares",
      "default_operator" : "AND"
    }
  },
  "sort" : {
    "rating" : "desc",
    "symbol" : "asc"
  }
}

```

<http://localhost:9200/cf etf/ search>

POST

HEADER

Content-type application/json

BODY

```

{
  "query" : {
    "query_string" : {
      "query" : "fund_name:ishares edge global",
      "default_operator" : "AND"
    }
  },
  "sort" : { "rating" : "desc", "morningstar_category" : "desc"
}
}

```

Searching using the query string:

<http://localhost:9200/cf etf/ search>

POST

HEADER

Content-type application/json

BODY

```

{
  "query" : {
    "query_string" : {
      "query" : "description:emerging non-diversified high-yield gobal division",
      "analyzer" : "description_analyzer"
    }
  }
}

```

Searching using the query string (return events which has all keywords in the search):

...

http://localhost:9200/cf_etf/_search

POST

...

HEADER

Content-type application/json

...

BODY

```
{
  "query" : {
    "query_string" : {
      "query" : "description:emerging non-diversified high-yield gobal division",
      "analyzer" : "description_analyzer",
      "default_operator" : "AND"
    }
  }
}
```

Searching using the query and rescore:

...

http://localhost:9200/cf_etf/_search

POST

...

HEADER

Content-type application/json

...

BODY

```
{
  "query" : {
    "query_string" : {
      "query" : "description:emerging non-diversified high-yield global dividend",
      "analyzer" : "description_analyzer"
    }
  },
  "rescore" : {
    "window_size" : 500,
    "query" : {
      "rescore_query" : {
        "query_string" : {
          "query" : "description:global",
          "analyzer" : "description_analyzer"
        }
      },
      "query_weight" : 0.1,
      "rescore_query_weight" : 0.9
    }
  }
}
```

Searching using the query and highlighting:

...

http://localhost:9200/cf_etf/_search

POST

...

HEADER

Content-type application/json

...

BODY

```
{
  "query" : {
    "bool" : {
      "should" : [
        {
          "query_string" : {
            "_name" : "Asia",
            "query" : "description:asia",
            "analyzer" : "description_analyzer"
          }
        },
        {
          "query_string" : {
            "_name" : "usa",
            "query" : "description:usa",
            "analyzer" : "description_analyzer"
          }
        },
        {
          "query_string" : {
            "_name" : "europe",
            "query" : "description:europe",
            "analyzer" : "description_analyzer"
          }
        }
      ]
    }
  },
  "highlight" : {
    "fields" : {
      "description" : {}
    }
  }
}
```

Multiple queries simultaneous execution:

...

<http://localhost:9200/cf etf/ search>

POST

...

HEADER

Content-type application/json

...

BODY

```
{
  "query" : {
    "bool" : {
      "should" : [
        { "query_string" : {
            "_name" : "asia",
            "query" : "description:asia",
```

```

        "analyzer" : "description_analyzer"}},
    {"query_string" : {
        "_name" : "europe",
        "query" : "description:europe",
        "analyzer" : "description_analyzer"}}
    ]
}

}

```

NOTE: Rescore can do 3 things filtering, ordering,

Search using OR condition:

...

http://localhost:9200/cf_etf/_search

POST

...

HEADER

Content-type application/json

...

BODY

```

{
    "query" : {
        "match" : {
            "fund_name" : {
                "query" : "ishares global"
            }
        }
    }
}

```

Search using Match with AND condition:

...

http://localhost:9200/cf_etf/_search

POST

...

HEADER

Content-type application/json

...

BODY

```

{
    "query" : {
        "match" : {
            "fund_name" : {
                "query" : "ishares global",
                "operator" : "and"
            }
        }
    }
}

```

Search using Match with AND condition and fuzziness:

...

http://localhost:9200/cf_etf/_search

POST

...

HEADER

Content-type application/json

...

BODY

```
{
  "query" : {
    "match" : {
      "fund_name" : {
        "query" : "ishares global",
        "operator" : "and",
        "fuzziness" : 2
      }
    }
  }
}
```

DAY 4

Match Queries:

Search using must and script:

...

http://localhost:9200/cf_etf/_search

POST

...

HEADER

Content-type application/json

...

BODY

```
{
  "query" : {
    "bool" : {
      "must" : {"match_phrase_prefix" : {"fund_name" : "ishares edge"}},
      "filter" : {
        "script" : {
          "script" : {
            "lang" : "painless",
            "source" : "if (!doc['rating'].empty && doc['rating'].value >= 4) return true;else false;"
          }
        }
      }
    }
  }
}
```

Search using must_not and script:

...

http://localhost:9200/cf_etf/_search

POST

...

HEADER

Content-type application/json

...

BODY

```
{
  "query" : {
    "bool" : {
      "must" : {"match_phrase_prefix" : {"fund_name" : "ishares edge"}},
      "must_not" : {"match" : {"fund_name" : "global" }}
    }
  }
}
```

Search using boosting:

...

http://localhost:9200/cf_etf/_search

POST

...

HEADER

Content-type application/json

...

BODY

```
{
  "query" : {
    "boosting":{
      "positive" : {"match_phrase_prefix" : {"fund_name" : "ishares edge"}},
      "negative" : {"term" : {"category" : "Equity"}},
      "negative_boost" : 0.1
    }
  }
}
```

Search using functions:

...

http://localhost:9200/cf_etf/_search

POST

...

HEADER

Content-type application/json

...

BODY

```
{
  "query" : {
```



```

    "function_score" : {
      "query" : {"match_phrase" : {"fund_name" : "ishares edge"}},
      "functions" : [
        {"filter" : {"wildcard" : {"morningstar_category" : "*Stock"}}, "weight" : 2},
        {"filter" : {"wildcard" : {"morningstar_category" : "*Growth"}}, "weight" : 1.
5},
        {"filter" : {"wildcard" : {"morningstar_category" : "*Blend"}}, "weight" : 0.5
      ]
    }
  }
}

```

Rank Evaluation:

...

http://localhost:9200/cf_etf/_rank_eval

POST

...

HEADER

Content-type application/json

...

BODY

```

{
  "requests" : [
    {
      "id" : "Asia",
      "request" : {"query" : {"match" : {"description" : "asia"}}},
      "ratings" : []
    },
    {
      "id" : "Global",
      "request" : {"query" : {"match" : {"description" : "global"}}},
      "ratings" : []
    }
  ],
  "metric" : {
    "precision" : {
      "k" : 3,
      "relevant_rating_threshold" : 1
    }
  }
}

```

Suggesters:

Provides auto completion sort of behavior. Faster search mechanism. e.g. songs suggestion based on your search.

Global Suggester:

Term Suggester:

Field base Suggester:

Completion Suggester:

Context Suggester:

Geo Context: Based on latitude and longitude you can search

Suggester:

...

http://localhost:9200/cf_etf/_rank_eval

POST

...

HEADER

Content-type application/json

...

BODY

```
{
  "suggest" : {
    "text" : "asian stock",
    "suggester_1" : {
      "term" : {
        "field" : "fund_name"
      }
    },
    "suggester_2" : {
      "term" : {
        "field" : "description"
      }
    }
  }
}
```

Suggester:

...

http://localhost:9200/cf_etf/_rank_eval

POST

...

HEADER

Content-type application/json

...

BODY

```
{
  "suggest" : {
    "text" : "asian stock",
    "suggester_1" : {
      "term" : {
        "field" : "fund_name"
      }
    },
    "suggester_2" : {
      "term" : {
        "field" : "description"
      }
    }
  }
}
```

```
    }}  
}
```

```
---
```

```
...  
http://localhost:9200/cf\_etf/\_rank\_eval
```

```
POST
```

```
...
```

```
HEADER
```

```
Content-type application/json
```

```
...
```

```
BODY
```

```
{  
  "suggest" : {  
    "text" : "divide weigh index",  
    "suggester" : {  
      "phrase" : {  
        "field" : "description"  
      }  
    }  
  }  
}
```

Create index:

```
curl -XDELETE "http://localhost:9200/mybooks"
```

```
curl -XPUT "http://localhost:9200/mybooks" -H 'Content-Type: application/json' -d'
```

```
{  
  "mappings": {  
    "properties": {  
      "join_field": {  
        "type": "join",  
        "relations": {  
          "order": "item"  
        }  
      },  
      "position": {  
        "type": "integer",  
        "store": true  
      },  
      "uuid": {  
        "store": true,  
        "type": "keyword"  
      },  
      "date": {  
        "type": "date"  
      },  
      "quantity": {  
        "type": "integer"  
      }  
    }  
  }  
}
```

```

    },
    "price": {
      "type": "double"
    },
    "description": {
      "term_vector": "with_positions_offsets",
      "store": true,
      "type": "text"
    },
    "title": {
      "term_vector": "with_positions_offsets",
      "store": true,
      "type": "text",
      "fielddata": true,
      "fields": {
        "keyword": {
          "type": "keyword",
          "ignore_above": 256
        }
      }
    }
  }
}
}
}
}'

```

```

curl -XPOST "http://localhost:9200/_bulk?refresh" -H 'Content-Type: application/json' -d'
{"index":{"_index":"mybooks", "_id":"1"}}
{"uuid":"11111","position":1,"title":"Joe Tester","description":"Joe Testere nice guy","date":"2015-10-22","price":4.3,"quantity":50}
{"index":{"_index":"mybooks", "_id":"2"}}
{"uuid":"22222","position":2,"title":"Bill Baloney","description":"Bill Testere nice guy","date":"2016-06-12","price":5,"quantity":34}
{"index":{"_index":"mybooks", "_id":"3"}}
{"uuid":"33333","position":3,"title":"Bill Klingon","description":"Bill is not\n          nice guy","date":"2017-09-21","price":6,"quantity":33}
'

```

```

curl -XDELETE "http://localhost:9200/mybooks-join"

```

```

curl -XPUT "http://localhost:9200/mybooks-join" -H 'Content-Type: application/json' -d'
{
  "mappings": {
    "properties": {
      "join": {
        "type": "join",
        "relations": {
          "book": "author"
        }
      }
    },
    "position": {
      "type": "integer",
      "store": true
    }
  },

```

```
"uuid": {
  "store": true,
  "type": "keyword"
},
"date": {
  "type": "date"
},
"quantity": {
  "type": "integer"
},
"price": {
  "type": "double"
},
"rating": {
  "type": "double"
},
"description": {
  "term_vector": "with_positions_offsets",
  "store": true,
  "type": "text"
},
"title": {
  "term_vector": "with_positions_offsets",
  "store": true,
  "type": "text",
  "fielddata": true,
  "fields": {
    "keyword": {
      "type": "keyword",
      "ignore_above": 256
    }
  }
},
"name": {
  "type": "text",
  "fields": {
    "keyword": {
      "type": "keyword",
      "ignore_above": 256
    }
  }
},
"surname": {
  "type": "text",
  "fields": {
    "keyword": {
      "type": "keyword",
      "ignore_above": 256
    }
  }
},
"versions": {
  "type": "nested",
```

```

    "properties": {
      "color": {
        "type": "keyword"
      },
      "size": {
        "type": "integer"
      }
    }
  }
}
}
}
}

```

```

curl -XPOST "http://localhost:9200/_bulk?refresh" -H 'Content-Type: application/json' -d'
{"index":{"_index":"mybooks-join", "_id":"1"}}
{"uuid":"11111","position":1,"title":"Joe Tester","description":"Joe Testere nice guy","date":"2015-10-22","price":4.3,"quantity":50,"join":{"name": "book"}, "versions":[{"color":"yellow", "size":5}, {"color":"blue", "size":15}]}
{"index":{"_index":"mybooks-join", "_id":"a1", "routing":"1"}}
{"name":"Peter","surname":"Doyle","rating":4.5,"join":{"name": "author", "parent":"1"}}
{"index":{"_index":"mybooks-join", "_id":"a12", "routing":"1"}}
{"name":"Mark","surname":"Twain","rating":4.2,"join":{"name": "author", "parent":"1"}}
{"index":{"_index":"mybooks-join", "_id":"2"}}
{"uuid":"22222","position":2,"title":"Bill Baloney","description":"Bill Testere nice guy","date":"2016-06-12","price":5,"quantity":34,"join":{"name": "book"}, "versions":[{"color":"red", "size":2}, {"color":"blue", "size":10}]}
{"index":{"_index":"mybooks-join", "_id":"a2", "routing":"2"}}
{"name":"Agatha","surname":"Princeton","rating":2.1,"join":{"name": "author", "parent":"2"}}
{"index":{"_index":"mybooks-join", "_id":"3"}}
{"uuid":"33333","position":3,"title":"Bill Klingon","description":"Bill is not\n      nice guy","date":"2017-09-21","price":6,"quantity":33,"join":{"name": "book"}, "versions":[{"color":"red", "size":2}]}
{"index":{"_index":"mybooks-join", "_id":"a3", "routing":"3"}}
{"name":"Martin","surname":"Twisted","rating":3.2,"join":{"name": "author", "parent":"3"}}
,

```

```

curl -XPOST "http://localhost:9200/mybooks-join/_refresh"

```

```

curl -XDELETE 'http://localhost:9200/mygeo-index'
curl -XPUT "http://localhost:9200/mygeo-index" -H 'Content-Type: application/json' -d'
{
  "mappings": {
    "properties": {
      "pin": {
        "properties": {
          "location": {
            "type": "geo_point"
          }
        }
      }
    }
  }
}
}
}

```

```
curl -XPUT 'http://localhost:9200/mygeo-index/_doc/1' -H 'Content-Type: application/json' -d '{"pin": {"location": {"lat": 40.12, "lon": -71.34}}}'
curl -XPUT 'http://localhost:9200/mygeo-index/_doc/2' -H 'Content-Type: application/json' -d '{"pin": {"location": {"lat": 40.12, "lon": 71.34}}}'
curl -XPOST 'http://localhost:9200/mygeo-index/_refresh'
```

Execute the below in Dev tools in kibana (you can also perform using Postman)

#prefix

POST /mybooks/_search

```
{
  "query": {
    "prefix": {
      "uuid": "222"
    }
  }
}
```

#wildcard

POST /mybooks/_search

```
{
  "query": {
    "wildcard": {
      "uuid": "22?2*"
    }
  }
}
```

#regex

POST /mybooks/_search

```
{
  "query": {
    "regexp": {
      "description": {
        "value": "j.*",
        "flags": "INTERSECTION|COMPLEMENT|EMPTY"
      }
    }
  }
}
```

#span queries

POST /mybooks/_search

```
{
  "query": {
    "span_first": {
      "match": {
        "span_term": {
          "description": "joe"
        }
      }
    },
  },
}
```

```
    "end": 5
  }
}
}
```

POST /mybooks/_search

```
{
  "query": {
    "span_or": {
      "clauses": [
        {
          "span_term": {
            "description": "nice"
          }
        },
        {
          "span_term": {
            "description": "cool"
          }
        },
        {
          "span_term": {
            "description": "wonderful"
          }
        }
      ]
    }
  }
}
```

POST /mybooks/_search

```
{
  "query": {
    "span_multi": {
      "match": {
        "prefix": {
          "description": {
            "value": "jo"
          }
        }
      }
    }
  }
}
```


POST /mybooks/_search

```
{
  "query": {
    "span_near": {
      "clauses": [
        {
          "span_term": {
            "description": "nice"
          }
        },
        {
          "span_term": {
            "description": "joe"
          }
        },
        {
          "span_term": {
            "description": "guy"
          }
        }
      ],
      "slop": 3,
      "in_order": false
    }
  }
}
```

POST /mybooks/_search

```
{
  "query": {
    "span_not": {
      "include": {
        "span_term": {
          "description": "nice"
        }
      },
      "exclude": {
        "span_near": {
          "clauses": [
            {
              "span_term": {
                "description": "not"
              }
            },
            {
              "span_term": {
                "description": "nice"
              }
            }
          ]
        }
      }
    }
  }
}
```

```
    "slop": 1,
    "in_order": true
  }
}
}
}
```

POST /mybooks/_search

```
{
  "query": {
    "span_containing": {
      "little": {
        "span_term": {
          "description": "nice"
        }
      },
      "big": {
        "span_near": {
          "clauses": [
            {
              "span_term": {
                "description": "not"
              }
            },
            {
              "span_term": {
                "description": "guy"
              }
            }
          ],
          "slop": 5,
          "in_order": true
        }
      }
    }
  }
}
```

POST /mybooks/_search

```
{
  "query": {
    "span_within": {
      "little": {
        "span_term": {
          "description": "nice"
        }
      },
      "big": {
```

```

"span_near": {
  "clauses": [
    {
      "span_term": {
        "description": "not"
      }
    },
    {
      "span_term": {
        "description": "guy"
      }
    }
  ],
  "slop": 5,
  "in_order": true
}
}
}
}

```

#match query

POST /mybooks/_search

```

{
  "query": {
    "match": {
      "description": {
        "query": "nice guy",
        "operator": "and"
      }
    }
  }
}

```

POST /mybooks/_search

```

{
  "query": {
    "match_phrase": {
      "description": "nice guy"
    }
  }
}

```

POST /mybooks/_search

```

{
  "query": {
    "match_phrase_prefix": {
      "description": "nice gu"
    }
  }
}

```

```

}

POST /mybooks/_search
{
  "query": {
    "multi_match": {
      "fields": [
        "description",
        "name"
      ],
      "query": "Bill",
      "operator": "and"
    }
  }
}

```

#query string query

```

POST /mybooks/_search
{
  "query": {
    "query_string": {
      "query": "\"\"nice guy\" -description:not price:{ * TO 5 } \"\"",
      "fields": [
        "description^5"
      ],
      "default_operator": "and"
    }
  }
}

```

#simple query string query

```

POST /mybooks/_search
{
  "query": {
    "simple_query_string": {
      "query": "\"\"nice guy\" -not\"",
      "fields": [
        "description^5",
        "_all"
      ],
      "default_operator": "and"
    }
  }
}

```

range

```

POST /mybooks/_search
{
  "query": {
    "range": {
      "position": {

```

```

    "from": 3,
    "to": 4,
    "include_lower": true,
    "include_upper": false
  }
}
}
}

```

#common text

POST /mybooks/_search

```

{
  "query": {
    "common": {
      "description": {
        "query": "nice guy",
        "cutoff_frequency": 0.001
      }
    }
  }
}
}

```

#ids

POST /mybooks/_search

```

{
  "query": {
    "ids": {
      "type": "test-type",
      "values": [
        "1",
        "2",
        "3"
      ]
    }
  }
}
}

```

Ingestion:

Will occur before indexing. e.g. enriching the data before indexing it.

Each Processors will be sequentially executed. Elasticsearch adds transforms this to index.

Whenever you create ingest pipeline (using script or using Kibana) both of them stores in a cluster.

Enabled by default in 8 and after versions. Pipelines are executed in ingest node. Pipeline has to be attached to the index (you can do at id level or doc level).

In Kibana -> stack management -> ingest pipeline

Every pipeline will have processors.

In Dev tools as well we can create pipeline.

Creating pipeline:

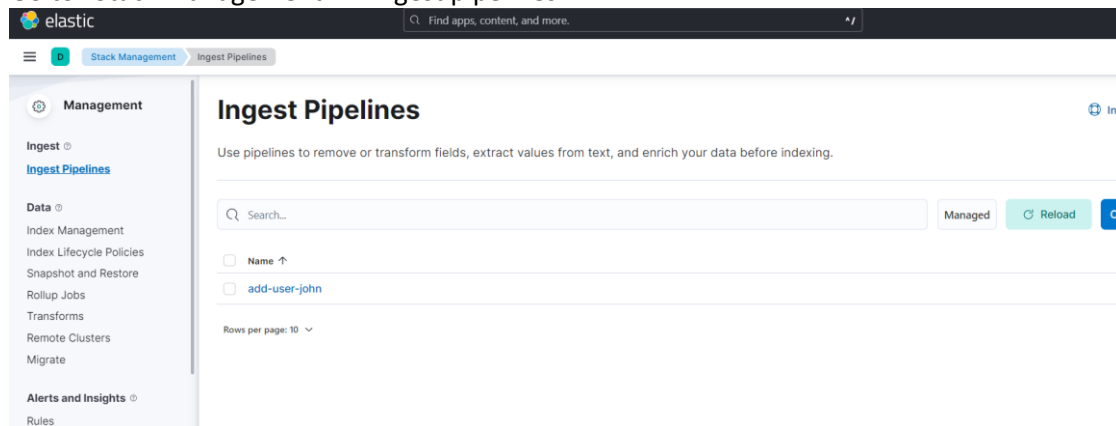
Got o Dev tools -> and execute below:

PUT /_ingest/pipeline/add-user-john

```
{
  "description": "Add user john field",
  "processors": [
    {
      "set" : {
        "field" : "user",
        "value": "john"
      }
    }
  ],
  "version" : 2
}
```

Check the pipeline created:

Go to “Stack management” -> ingest pipelines



PUT /my_index/_doc/my_id?pipeline=add-user-john

{}

GET /my_index/_doc/my_id

GET /_ingest/pipeline/add-user-john

GET /_ingest/pipeline/*

GET /_ingest/pipeline/add-*

DELETE /_ingest/pipeline/add-user-john

#Simulate an ingest pipeline

POST /_ingest/pipeline/_simulate

{

```

"pipeline": {
  "description": "Add user john field",
  "processors": [
    {
      "set": {
        "field": "user",
        "value": "john"
      }
    },
    {
      "set": {
        "field": "job",
        "value": 10
      }
    }
  ],
  "version": 1
},
"docs": [
  {
    "_index": "index",
    "_type": "type",
    "_id": "1",
    "_source": {
      "name": "docs1"
    }
  },
  {
    "_index": "index",
    "_type": "type",
    "_id": "2",
    "_source": {
      "name": "docs2"
    }
  }
]
}

```

Multiple processors can be created. You can have different step by step processes. We can do the simulation. We can perform different types of activities. Simulation is for testing purpose without creating any object to test the pipeline.

DAY 5

Activities performing with multiple processors:

POST /_ingest/pipeline/_simulate

```

{
  "pipeline": {
    "description": "Testing some build-processors",

```

```
"processors": [
  {
    "dot_expander": {
      "field": "extfield.innerfield"
    }
  },
  {
    "remove": {
      "field": "unwanted"
    }
  },
  {
    "trim": {
      "field": "message"
    }
  },
  {
    "set": {
      "field": "tokens",
      "value": "{{message}}"
    }
  },
  {
    "split": {
      "field": "tokens",
      "separator": "\\s+"
    }
  },
  {
    "sort": {
      "field": "tokens",
      "order": "desc"
    }
  },
  {
    "convert": {
      "field": "mynumbertext",
      "target_field": "mynumber",
      "type": "integer"
    }
  }
],
"docs": [
  {
    "_index": "index",
    "_type": "type",
    "_id": "1",
    "_source": {
      "extfield.innerfield": "booo",
      "unwanted": 32243,
      "message": "155.2.124.3 GET /index.html 15442 0.038",
      "mynumbertext": "3123"
    }
  }
]
```



```
}  
}  
]  
}
```

Grok filter:

POST /_ingest/pipeline/_simulate

```
{  
  "pipeline": {  
    "description": "Testing grok pattern",  
    "processors": [  
      {  
        "grok": {  
          "field": "message",  
          "patterns": [  
            "%{IP:client} %{WORD:method} %{URIPATHPARAM:request} %{NUMBER:bytes} %{NUMBER:duration}"  
          ]  
        }  
      }  
    ]  
  },  
  "docs": [  
    {  
      "_index": "index",  
      "_type": "type",  
      "_id": "1",  
      "_source": {  
        "message": "155.2.124.3 GET /index.html 15442 0.038"  
      }  
    }  
  ]  
}
```

POST /_ingest/pipeline/_simulate

```
{  
  "pipeline": {  
    "description": "custom grok pattern",  
    "processors": [  
      {  
        "grok": {  
          "field": "message",  
          "patterns": [  
            "my favorite color is %{COLOR:color}"  
          ],  
          "pattern_definitions": {  
            "COLOR": "RED|GREEN|BLUE"  
          }  
        }  
      }  
    ]  
  }  
}
```

```

},
"docs": [
  {
    "_source": {
      "message": "my favorite color is RED"
    }
  },
  {
    "_source": {
      "message": "happy fail!!"
    }
  }
]
}

```

Encode in pipeline:

#Using the ingest attachment plugin

PUT /_ingest/pipeline/attachment

```

{
  "description": "Extract data from an attachment via Tika",
  "processors": [
    {
      "attachment": {
        "field": "data"
      }
    }
  ],
  "version": 1
}

```

PUT /my_index/_doc/my_id1?pipeline=attachment

```

{
  "data":
  "SWQsIE1vZGVsLCBZZWFyLCBTY3JlZW5TaXplLCBSQU0sIEhERCwgVywgRCwgSCwgV2VpZ2h0CjEsTWJfQm9vayBQc
m8sMjAxNSwiMTVcliIsMTZHQiw1MTJHQBtU0QsMTMuNzUsOS40OCwwLjYxLDQuMDIKMixNYWNCb29rLDIwMTY
sljEYXCIiLDhHQiwyNTZHQiBTU0QsMTEuMDQsNy43NCwwLjUyLDIuMDMKMyxNYWNCb29rIEFpciwyMDE2LCIxMy4
zXCIiLDhHQiwxMjhHQiBTU0QsMTIuOCw4Ljk0LDAuNjgsMi45Ngo0LGINYWMsMjAxNywiMjdcliIsNjRHQiwxVElgU1N
ELDI1LjYsOC4wLDIwLjMsMjAuOA=="
}

```

GET /my_index/_doc/my_id1

Geo IP concept in pipeline:

#Using the ingest GeoIP plugin

PUT /_ingest/pipeline/geoip

```

{
  "description": "Extract geopoint from an IP",

```

```

"processors": [
  {
    "geoip": {
      "field": "ip"
    }
  }
],
"version": 1
}

```

PUT /my_index/_doc/my_ip?pipeline=geoip

```

{
  "ip" : "46.51.224.0"
}

```

GET /my_index/_doc/my_ip

LogStash

Is an independent product. You have different categories of activities that can be performed. Logs stash configuration using logstash.conf.

Edit the logstash.conf with the below:

```

input {
  stdin {}
}

output {
  elasticsearch {
    hosts => ["localhost:9200"]
    index => "oracle_stash"
  }
}

```

Now go to directory: C:\Users\abraturi\Desktop\ELK\logstash-8.8.0\bin

Once done execute the below command: .\logstash.bat -f .\logstash.conf

Now provide data by simply typing and hit enter:

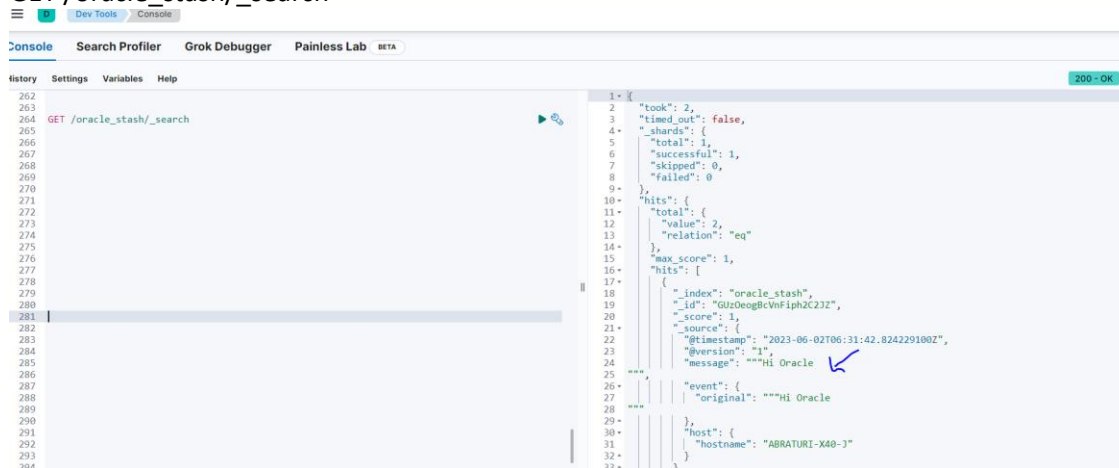
```

[2023-06-02T12:00:39,316][INFO ][logstash.outputs.elasticsearch][main] Not eligible for data streams because config contains one or more settings that are not
th data streams: {"index"=>"oracle_stash"}
[2023-06-02T12:00:39,326][INFO ][logstash.outputs.elasticsearch][main] Data streams auto configuration ('data_stream' => auto' or unset) resolved to 'false'
[2023-06-02T12:00:39,328][WARN ][logstash.outputs.elasticsearch][main] Elasticsearch Output configured with 'ecs_compatibility => v8', which resolved to an UNREL
ew of version 8.0.0 of the Elastic Common Schema. Once ECS v8 and an updated release of this plugin are publicly available, you will need to update this plugin t
his warning.
[2023-06-02T12:00:39,379][INFO ][logstash.javapipeline ][main] Starting pipeline {:pipeline_id=>"main", "pipeline.workers"=>8, "pipeline.batch.size"=>125, "pi
h.delay"=>50, "pipeline.max_inflight"=>1000, "pipeline.sources"=>["C:/Users/abraturi/Desktop/ELK/logstash-8.8.0/logstash.conf"]}, :thread=>"#<Thread:0x73b792cd@C:
turi/Desktop/ELK/logstash-8.8.0/logstash-core/lib/logstash/java_pipeline.rb:134 run>"}
[2023-06-02T12:00:39,396][INFO ][logstash.outputs.elasticsearch][main] Using a default mapping template {:es_version=>8, :ecs_compatibility=>v8}
[2023-06-02T12:00:40,235][INFO ][logstash.javapipeline ][main] Pipeline Java execution initialization time {"seconds"=>0.83}
[2023-06-02T12:00:40,342][INFO ][logstash.javapipeline ][main] Pipeline started {"pipeline.id"=>"main"}
The stdin plugin is now waiting for input:
[2023-06-02T12:00:40,358][INFO ][logstash.agent ][main] Pipelines running {:count=>1, :running_pipelines=>[:main], :non_running_pipelines=>[]}
Hi Oracle
dsfjf

```

Check the data in kibana using Dev tools:

GET /oracle_stash/_search



Filebeat configuration:

Logstash and beat has to be on same machine.

Install file beat and put the filebeat-example.yml file there

Now under logstash put the filebeat-demo.conf file.

Also put the logstash-dataset under the same path

https://github.com/vishymails/ELK_RAW/tree/main

Logstash will collect the data and send to elasticsearch and elasticsearch will store in nosql database in the form of indexes.

Using python to connect to Elasticsearch:

Install Anaconda Navigator

Open Jupyter lab from the Anaconda Navigator

Create ipynb file with below code

```
!pip install elasticsearch
```

```
!pip install pandas
```

```
print("ELK and python client example")
```

```
try:
```

```
    import sys
```

```
    import os
```

```
    import elasticsearch
```

```
    from elasticsearch import Elasticsearch
```

```
    import pandas as pd
```

```
    print("all libraries are imported")
```

```
except Exception as e:
```

```
    print("some modules are not imported",e)
```

```
def connect_elasticsearch():
```

```

es=None
es=Elasticsearch('http://localhost:9200')

if es.ping():
    print("Elastic server connected")

else:
    print("Elastic server not connected")

return es

```

```
es = connect_elasticsearch()
```

```

[12]: !pip install elasticsearch
      #!pip install pandas

Requirement already satisfied: elasticsearch in c:\users\abratur\anaconda3\lib\site-packages (8.8.0)
Requirement already satisfied: elastic-transport<9,>=8 in c:\users\abratur\anaconda3\lib\site-packages (from elasticsearch) (8.4.0)
Requirement already satisfied: urllib3<2,>=1.26.2 in c:\users\abratur\anaconda3\lib\site-packages (from elastic-transport<9,>=8->elasticsearch) (1.26.14)
Requirement already satisfied: certifi in c:\users\abratur\anaconda3\lib\site-packages (from elastic-transport<9,>=8->elasticsearch) (2022.12.7)

[3]: print("ELK and python client example")

ELK and python client example

[5]: try:
      import sys
      import os
      import elasticsearch
      from elasticsearch import Elasticsearch
      import pandas as pd

      print("all libraries are imported")

    except Exception as e:
      print("some modules are not imported",e)

all libraries are imported

[9]: def connect_elasticsearch():
      es=None
      es=Elasticsearch('http://localhost:9200')

      if es.ping():
          print("Elastic server connected")

      else:
          print("Elastic server not connected")

      return es

[10]: es = connect_elasticsearch()

Elastic server connected

```

```

es.indices.create(index="python-index", ignore=400)
es.indices.create(index="python-index1", ignore=400)

```

```

res=es.indices.get_alias(index="*")
for Name in res:
    print(Name)

```

```

from datetime import datetime
doc={
    'author' : 'ABhishek Raturi',
    'text' : 'Elastic Search Example',
    'customer' : 'Oracle India',
    'session_date' : datetime.now()
}
resp=es.index(index="python-index",id=2,document=doc)

```

```
print(resp['result'])
resp = es.get(index='python-index', id=1)
print(resp['_source'])
```

```
from datetime import datetime
doc={
    'author' : 'ABhishek Raturi',
    'text' : 'Elastic Search Example2',
    'customer' : 'Oracle India',
    'session_date' : datetime.now()
}
resp=es.index(index="python-index",id=2,document=doc)
print(resp['result'])
```

```
es.delete(index='python-index', id=1)
```